

Testing the Public Goods Theory of Alliances

Joshua Alley*

March 5, 2019

Abstract

Olson and Zeckhauser (1966)’s public goods model of alliances exerts a large influence on scholarship and policy, but it has not subjected to careful empirical scrutiny. Prior tests of the public goods argument suffer from a mix of identification and generalizability problems. In this paper, I examine two predictions from the public goods model. First, I test whether state size modifies the impact of increasing allied capability on growth in military spending. Then I estimate the association between treaty contribution and growth in military spending for 285 offensive and defensive alliances. Neither design supports the expectations of the public goods model. Therefore, the argument that alliances provide a public good and generate free-riding should be treated with more skepticism.

1 Introduction

Olson and Zeckhauser (1966) argue that international alliances are subject to a collective action problem. In their model, security from an alliance is a public good, which allows smaller alliance participants “free ride” on the contributions of larger members. Free-riding is reflected in lower defense burdens— states allocate a lower share of their resources to the military. This view of alliances and defense effort is quite influential.

*Graduate Student, Department of Political Science, Texas A&M University.

As of February 2019, Olson and Zeckhauser's article has 1624 citations, making it the best-cited work on alliance participation and military spending by a wide margin. But this "economic theory of alliances" is even more influential outside of the academy. Policy and popular discussions of alliances are full of references to free-riding.

Despite the influence of the public goods model of alliances, we do not know how accurate this argument is. Previous empirical tests of Olson and Zeckhauser's predictions have identification problems and lack generalizability. In this paper, I offer the first comprehensive test of the public goods theory of alliances, and find little support for its predictions about state size and military spending. My findings have important implications for the way scholars and policymakers discuss the connection between alliances and defense expenditures.

US policymakers make frequent use of free-riding and disproportionate contributions to the common good to criticize lackluster allied defense expenditures. For example, Barack Obama complained in 2016 that "Free riders aggravate me" and US allies "have to pay your fair share." Donald Trump has implied the US will not protect allies who spend too little on defense. Such exhortations go back to the Eisenhower administration (Lanoszka, 2015).

In US alliance politics, the North Atlantic Treaty Organization (NATO) is the epicenter of free-riding discussions. Following Olson and Zeckhauser's emphasis on the defense burden, accusations of free-riding emphasize how NATO allies spend a lower share of their income on defense than the US. Pundits and policy makers now employ spending 2% of GDP on defense as a threshold for adequate effort. NATO members made a formal commitment to meet this spending target in 2014.¹

Failure of NATO members to bear an adequate defense burden is then seen as evidence of free-riding. US policymakers assert low allied defense burdens are an inadequate contribution to collective defense. European allies reject this characterization of their defense spending, arguing they make other crucial contributions to common defense. Academic studies echo this debate about

¹In 2017 only four NATO members besides the US met the 2% threshold.

free-riding. Scholars have spent decades arguing over whether defense burdens of NATO members reflect free-riding (Sandler and Forbes, 1980; Palmer, 1990; Gates and Terasawa, 1992; Sandler and Hartley, 2001; Lanoszka, 2015; Plümper and Neumayer, 2015).

Olson and Zeckhauser's argument shaped a generation of policy and scholarly debate about alliances and defense spending. But 52 years after the publication of "An Economic Theory of Alliances," we have limited evidence for or against their empirical predictions. Therefore, it is still unclear whether the public goods approach to alliances is a good model of alliance politics. The lack of reliable evidence is the consequence of two problems.

First, many empirical estimates of defense burdens within alliances are unidentified. Second, almost every test examines NATO. Olson and Zeckhauser's measures of the defense burden and state size create an identification problem. They estimated correlations between military spending as a share of national income and national income. While this approach is widely emulated, it places GDP on both sides of the equation. Therefore, most empirical models of the public goods theory are unidentified.

In Olson and Zeckhauser's test, changes in GDP impact the independent and dependent variable. There is a deterministic component in the relationship between GDP and the defense burden—changes in GDP shift the defense burden.² Moreover, ratio dependent variables such as military expenditures as a share of GDP often generate spurious results (Kronmal, 1993).

One noteworthy paper addresses the identification problem, but may not produce generalizable findings. Plümper and Neumayer (2015) examine how changes in spending by NATO members respond to US and Soviet spending. In their framework, unresponsiveness to US and Soviet military spending is evidence of free riding by non-US NATO members. They find no correlation between NATO member size and free-riding, which they argue contradicts Olson and Zeckhauser.³

²The only exception occurs when military spending changes in such a way that defense spending's share of GDP is unchanged. Such changes are possible, but highly unlikely.

³However, they do not include the United States in their sample, which omits a crucial data point for testing the size argument. Given the focus on responses to US spending, this decision is understandable.

Plümper and Neumayer (2015)'s focus on NATO reflects the other major research design shortcoming—a lack of generalizability. Most studies focus on NATO and most evidence suggests NATO members spend less on the military thanks to allied capability (George and Sandler, 2017). But NATO is an unusually capable and durable treaty, making it a difficult case to draw general conclusions from.

We do not know whether the public goods theory of alliances applies to treaties besides NATO. My survey of the literature on alliance participation and military spending found six tests of the public goods theory of alliances outside of NATO. All six of those studies include GDP in the independent and dependent variable. These tests of the public goods theory outside of NATO are unreliable.

In summary, the predictions of the economic theory of alliances have not been subjected to a general and well-identified test. The best identified test may not be generalizable and other empirical models are unidentified. These research design limitations have two important consequences.

First, they hinder the accumulation of knowledge. Without a valid and comprehensive test, the explanatory power of the public goods theory of alliances is unclear. Thus, it is hard to determine whether other theories should build on this model. Arguments that treat security from an alliance as a public good could rest on shaky foundations.

Other scholars have questioned and revised the underlying logic of the public goods theory of alliances (Palmer, 1990; Sandler and Hartley, 2001). But theoretical revisions are premature without evidence a parsimonious public goods model is inappropriate. We should not rule out an intuitive public goods theory without evidence it makes inaccurate predictions.⁴

The other crucial consequence of poor tests is that using free-riding to describe reduced defense spending by allied states may have little empirical justification. Free-riding is inseparable from public goods and collective action. Popular and policy discussions of alliance politics implicitly depend on the public goods model of alliances. Policy efforts to address lower military spending

⁴Findings of no correlation between GDP and the defense burden have the same identification problem.

could fail if they are based on an inaccurate model of alliance politics.

Therefore, establishing the empirical validity of the public goods theory of alliances is crucial for academic theory and policy debates. In this paper, I address the identification and generalizability issues in two novel tests the public goods model of alliances. Both research designs examine how alliance participant size shapes growth in military spending.

The first design uses a conventional panel data model with an aggregate measure of alliance participation. I interact state size and changes in allied spending examine whether states reduce military expenditures as a result of increasing allied spending, and if reductions diminish with greater state size. The second design employs a Bayesian model to estimate the association between treaty contribution and growth in military spending for individual alliances. Neither approach finds evidence to match Olson and Zeckhauser's prediction that greater size requires more military spending.

The paper proceeds as follows. First, I summarize the public goods theory of alliances and two predictions of its logic. Then, I describe the results of a panel-data test of the first prediction. The third section tests the second prediction with a Bayesian multilevel model. The final section assesses aggregate support for the public goods logic, as well as implications for scholarship and policy.

2 The Public Goods Theory of Alliances

Why might alliances suffer from a public goods problem? The aggregate military capability of an alliance provides security for members. Because a treaty cannot exclude members without undermining its purpose, security is a public good.

Treaty members receive security regardless of their individual contribution. Therefore states have incentives to rely on other members to contribute military capability to the treaty. Olson and Zeckhauser expect larger members of the alliance will bear a higher defense burden, because these

states value defense from the alliance more. Smaller alliance members rely on the contributions of larger partners for security and bear a lower defense burden. As a result, smaller states exploit larger alliance participants by free riding.

There are two implications of this argument. One, large and small states should respond differently to changes in allied military spending. If Olson and Zeckhauser's argument is correct, smaller states should decrease military spending in response to greater allied military spending. Larger states will be less responsive to increased allied spending. This implies a conditional relationship between allied spending, state size, and defense expenditures.

HYPOTHESIS 1: The marginal effect of allied military spending on growth in state-year military spending will be negative for small states, and increasing in state size.

The other prediction addresses individual alliance treaties. If larger states bear a disproportionate share of the alliance burden, then within each alliance, increasing contribution will require extra military spending. States that contribute more to an alliance will increase military expenditures to maintain security from the treaty.

HYPOTHESIS 2: For individual alliance treaties, increasing treaty contribution will be positively associated with annual growth in military spending.

Treaties with a positive correlation between contribution and military spending conform to Olson and Zeckhauser's expectations. No correlation indicates larger members are not increasing military expenditures. A negative correlation between treaty contribution and military spending implies larger members decrease spending.

I now test these two predictions using a corresponding research design in a sample of all states except microstates from 1816 to 2007.⁵ State-year observations are the unit of analysis in both tests. The next section examines Hypothesis 1 by interacting changes in total allied spending and GDP. This approach aggregates multiple treaties into a single measure of alliance participation.

⁵Including states with small GDP values makes estimating the interaction in Hypothesis 1 difficult. These observations create a large part of the modifying variable with no support for interactive estimates.

3 Testing Hypothesis 1

To test Hypothesis 1, I employ conventional measures of state size and alliance participation. Olson and Zeckhauser use GDP as a proxy for state size, so I constructed a measure of GDP using data from the Maddison Project (Bolt et al., 2018). I then take the natural log of GDP to pull in outliers. I measure total allied capability as the military spending of all defensive or offensive alliance partners of a state. All alliance membership data comes from the Alliance Treaty Obligations and Provisions (ATOP) data (Leeds et al., 2002).

Growth in the natural log of military spending is the dependent variable (Singer, 1988). The log-level of military spending is non-stationary. Differenced military spending has increasing variance over time, as budgets expand and generate larger changes. Modeling the DV in levels or changes with a control for the level might create spurious inferences (Granger and Newbold, 1974).

Using growth in military spending as the dependent variable benchmarks changes to budget size. Change in spending occurs with reference to prior expenditures. This mitigates the risk of non-stationarity in changes or levels of military spending generating spurious inferences. Growth in spending is equal to:

$$\text{Growth Military Spending} = \frac{\text{Mil. Ex.}_t - \text{Mil. Ex.}_{t-1}}{\text{Mil. Ex.}_{t-1}} = \frac{\Delta \text{Mil. Ex.}}{\text{Mil. Ex.}_{t-1}} \quad (1)$$

Because Hypothesis 1 implies a conditional relationship, I interact allied spending and GDP in the following regression model:

$$y = \beta_1 + \beta_2 \ln(\text{GDP}) + \beta_3 \text{Ally Spend} + \beta_4 \ln(\text{GDP}) \times \text{Ally Spend} + \beta \mathbf{X} + \epsilon \quad (2)$$

β_2 and β_3 are the constituent terms of the interaction and \mathbf{X} is a matrix of control variables with coefficient vector β . Hypothesis 1 implies β_2 should be positive, and β_3 should be negative. The interaction term β_4 should be positive.

I estimate this model with a robust regression estimator. Growth in military spending has heavy-tailed residuals in a regression, so OLS is inefficient (Rainey and Baissa, 2018). Given heavy-tailed residuals, non-linear estimators such as robust regression can outperform OLS. Robust regression re-weights observations using the size of the residual, making it less sensitive to large deviations than OLS.

X contains correlates of alliance participation and military spending. I control for international war (Reiter, Stam and Horowitz, 2016), civil war participation (Sarkees and Wayman, 2010), and a count of annual MIDs (Gibler, Miller and Little, 2016). Other controls average salient alliance characteristics, including democratic membership (DiGiuseppe and Poast, 2016) and number of members across a state's alliance treaties. Last, I control for regime type, external threat (Leeds and Savun, 2007), and Cold War years.

3.1 Results

Estimates from this model do not match the expectations of Hypothesis 1. Table 1 summarizes results from the robust regression estimator. The log of GDP and changes in allied spending have no effect on growth in military spending. The interaction between changes in allied spending and GDP is also statistically insignificant.

I cannot infer the absence of an interaction from the coefficients alone (Brambor, Clark and Golder, 2006). Figure 1 plots the marginal effect of allied military spending across the range of GDP. The marginal effect of allied military spending can only be distinguished from zero for positive values. There is no conditional relationship between state size and changes in allied spending.

These results hold if I replace robust regression with OLS or employ kernel estimation of the interactive relationship (Hainmueller, Mummolo and Xu, 2019).⁶ I also estimated a model using states with an alliance as the estimation sample with Heckman two-stage estimator to address non-random selection into alliances. In the aggregate, state size does not modify the impact of allied

⁶See the appendix.

| | Growth Military Spending |
|---|--------------------------|
| Change Allied Spending | −0.026 (0.047) |
| ln(GDP) | 0.0003 (0.001) |
| Change Allied Spending \times ln(GDP) | 0.002 (0.002) |
| Avg. Alliance Size | 0.0003* (0.0002) |
| Avg Alliance Democracy | −0.0003 (0.001) |
| International War | 0.097*** (0.010) |
| Civil War Participant | 0.001 (0.008) |
| POLITY | −0.0001 (0.0004) |
| External Threat | 0.043*** (0.011) |
| Cold War | 0.047*** (0.004) |
| Constant | 0.026 (0.030) |
| Observations | 9,139 |
| Residual Std. Error | 0.157 (df = 9128) |

Note: *p<0.1; **p<0.05; ***p<0.01

Table 1: Robust Regression Estimates: Association of interaction of allied spending and GDP with growth in military spending 1816-2007.

Marginal Effect of Changes in Allied Spending on Military Spending




Figure 1: Average Marginal Effect of increasing allied military spending on a state's military spending, across the range of GDP. Rug plot on x-axis summarizes the distribution of GDP.

capability on growth in military spending.

However, it is possible that aggregating alliance treaties obscures collective action dynamics in individual treaties. Olson and Zeckhauser focus on relative size, so states with small GDP can make large alliance contributions if their partner is even smaller. The next section examines how often greater treaty contribution increases military expenditures.

4 Testing Hypothesis 2

Testing Hypothesis 2 requires estimating the association between alliance contribution and military spending for each alliance. For each of the 285 alliances that promise military support,⁷ I estimate a parameter measuring the impact of increasing alliance contribution on military expenditures. Bayesian estimation regularizes estimates with many parameters, so I fit the following model using STAN (Carpenter et al., 2016).

The full model starts with state-year growth in military spending y_{it} . I model the DV with a t-distribution to account for heavy tails. ν is the degrees of freedom parameter— as ν increases, the t distribution becomes more like a normal distribution.

$$y_{it} \sim student_t(\nu, \mu, \sigma) \quad (3)$$

Most of this model follows standard panel data designs. The expected value of the outcome μ_{it} depends on a constant α , state and year varying intercepts α^{st} and α^{yr} and control variables $X_{it}\beta$. In this specification, I include all controls besides the alliance portfolio averages from the first robust regression in X .

$$\mu_{it} = \alpha + \alpha^{st} + \alpha^{yr} + \mathbf{X}_{it}\beta + \mathbf{Z}_{it}\gamma \quad (4)$$

⁷ATOP offensive and defensive treaties.

The $\mathbf{Z}_{it}\gamma$ term captures the impact of multiple alliances. \mathbf{Z} is a matrix of state participation in alliances—columns are alliances, rows are state-year observations. If a state is not in an alliance, the corresponding element of the matrix is equal to zero. If a state is part of an alliance, the corresponding element of the matrix is equal to a state’s military spending as a share of total alliance expenditures. The alliance contribution elements of the matrix range from 0 to 1, because some states have no military spending.⁸

\mathbf{Z} is a quasi-spatial approach to capturing the impact of participation in multiple alliances. According to Olson and Zeckhauser, greater treaty contribution should increase demand for military spending. Lower contribution will lead to reduced growth in spending, because these states can free-ride on their partners. The correlation between treaty contribution and growth in spending for each alliance is expressed in the γ parameters.

γ is a vector of 285 alliance-specific parameters. Because \mathbf{Z} contains alliance contribution, these coefficients estimate the association between treaty contribution and military spending. When a state is not in an alliance, the corresponding γ is multiplied by zero, and has no impact. Alliance participation only affects military expenditures if a state contributes.

Each alliance a state is a member of has a separate impact on military spending. Thus, every alliance estimate holds the impact of other treaties constant. A positive γ implies that as contribution to the alliance increases, members spend more, as Hypothesis 2 predicts.

4.1 Results

If the public goods theory of alliances is correct, we should observe many positive γ parameters. Because I employed Bayesian modeling, each γ has a posterior distribution.⁹ I focus interpretation on the posterior mean and 90% credible intervals.¹⁰ The posterior mean is the expected value of γ , while the credible intervals capture uncertainty around that estimate.

⁸Costa Rica is the best-known example of a state with no military spending, but these cases are highly unusual.

⁹See the appendix for a full summary of priors, convergence and model fit.

¹⁰I use 90% credible intervals because inferences around 95% intervals can be less stable.




Figure 2: Posterior mean of association between alliance contribution and military spending in 285 defensive and offensive alliances from 1816 to 2007.

Figure 2 summarizes the posterior mean of the 285 alliance coefficients. Most means are close to zero. There are large positive values, but also many large negative values.

147 of 285 alliances have a positive posterior mean. 138 of 285 alliances have a negative posterior mean. The large number of negative values provides some evidence against Hypothesis 2. However, the posterior means do not convey whether the impact of increasing alliance contribution can be distinguished from zero.

Figure 3 plots the γ parameter for each alliance against the start year of the treaty. Points mark the posterior mean. The error bars encapsulate the 90% credible interval.¹¹




Figure 3: Estimated association between alliance contribution and defense spending in 285 defensive and offensive alliances from 1816 to 2007. Points represent the posterior mean and the error bars cover the 90% credible interval.

¹¹The credible interval falls between the 5% and 95% quantiles of the posterior.

Most posterior means are close to zero, so 251 credible intervals include zero. There are 16 alliances where the credible interval includes only positive values, as the public goods model predicts. That leaves 13 treaties where the credible interval covers only negative values.

16 of 285, or 6%, is a dismal prediction success rate for Hypothesis 2. There is no association between treaty contribution and growth in military spending in most alliances. The 13 negative γ parameters imply larger members of those alliances *lower growth in military spending*, which the public goods theory cannot explain.

Alliances with a non-zero γ provide additional information. I examine these cases in more detail to assess whether collective action is the main issue. In Figure 4, I plot the posterior means of γ for all 29 alliances. Several of these alliances have a large expected impact on military spending.

The largest positive association between treaty contribution and defense spending growth is the 1961 Defense Pact of the African and Malagasy Union (ATOPID 3455). The second largest is an 1866 offensive pact between Italy and Prussia (ATOPID 1265). Third largest is an Inter-American Hemispheric defense pact during World War II (ATOPID 3010).

None of these three cases are obvious examples of free-riding. The Defense Pact of the African and Malagasy Union included most Francophone countries. The Union changed its functions to focus on economic affairs in 1964, dissolving the defense pact. Because the members of this treaty were geographically disparate and depended on France the defense functions of the treaty were weak from the start.

The Italy-Prussia agreement and the Inter-American Hemispheric defense pact formed immediately before and during war, respectively. These treaties are a special cases in Olson and Zeckhauser's model where defense spending is a superior good— increases in income are spent on defense. Larger members spent more because they used the treaty to expand their war-fighting efforts.

Alliances with large negative γ parameters have little in common with a public goods model. The largest negative association is an 1870 alliance between Prussia and the UK, signed during




Figure 4: Bar plot of γ for alliances where treaty contribution is positively or negatively correlated with military spending. The Y axis is the ATOP project's alliance identifier.

the Franco-Prussian war (ATOPID 1295). The second largest negative γ is a June 1939 alliance between France and Turkey aimed at securing the Mediterranean (ATOPID 2465). Last, the third most negative γ is a 1944 agreement between the US and Portugal, where the Portuguese agreed to help in the Pacific theater.

These three negative cases suggest states with resource constraints can use alliances to secure their interests and limit spending growth. Without a treaty with Prussia to guarantee Belgian neutrality, the UK would have spent more in 1870. France used the 1939 treaty with Turkey for similar purposes. The US formed the 1944 treaty with Portugal to further weaken Japan without exerting more military effort. All three alliances structure specific exchanges among members, rather than provide a public good.

NATO is the best example of collective defense, where other results suggest public goods dynamics are present. The estimated γ for NATO offers no support for the public goods theory of alliances, however. The posterior mean is -0.07 , and the credible interval ranges from $-.30$ to $.14$. Greater contribution to NATO is unassociated with growth in military spending. This finding corroborates the results of Plümper and Neumayer (2015).

A few cases may be subject to free-riding. The 2006 African Union Non-Aggression and Common Defense Pact (ATOPID 5055) has a positive γ that might be explained by free-riding. China and North Korea's 1961 alliance (ATOPID 3445) is another possible case— North Korea might have spent even more on the military without the Chinese alliance. Most of the other positive, non-zero γ parameters are from treaties formed near international conflict.

This second set of estimates also provides little support for the public goods theory of alliances. In 88% of alliances, there is no association between treaty contribution and military spending. Even cases where treaty contribution is positively correlated with spending may not reflect collective action problems.

5 Conclusion

Taking the results from testing Hypothesis 1 and 2 together, there is little evidence for the predictions of the public goods theory of alliances. There is no conditional relationship between GDP and changes in allied spending. Moreover, there is a clear positive association between treaty contribution and military spending in only 6% of alliances.

These findings should increase our skepticism about the public goods model of alliances. Although Olson and Zeckhauser's model is simple and intuitive, it lacks explanatory power. Better identified empirical models and estimates from multiple treaties show little evidence of collective action problems.

My results reinforce theoretical skepticism about the public goods model. Critiques by Palmer (1990) and Sandler and Hartley (2001) merit further attention. Constructing parsimonious models of alliances and defense effort is the next challenge for scholars. Components of the public goods model and these alternatives may provide a useful starting point.

Some of the results raise additional questions for theoretical work. Why do we observe 14 alliances where increasing alliance contributions are associated with reduced spending? Why is the effect of greater allied spending mostly positive in Model 1? Derivatives of the public goods theory will struggle to answer these questions.

Perhaps large and small states use alliances for distinct purposes, so members reap specific gains from treaty participation (Morrow, 1991). Small states might use alliances to defend their homeland. Large states could be more focused on using alliances to defend other states and expand their influence abroad.

Scholars and policymakers should reassess the way they use "free-riding" to describe alliance politics. Free-riding is inextricable from a public goods understanding of alliances. But if key predictions of the public goods model have little empirical support, free-riding is an inaccurate description of reduced defense effort by alliance participants. Charges of free-riding are normati-

vely loaded and may mask exchange between alliance members (Lanoszka, 2015).

We should not completely abandon the public goods model in international politics. Other international organizations could suffer from collective action problems. International environmental regimes and climate change are one noteworthy example. Using alliances as exemplars of collective action problems in other international organizations is inappropriate, however.

Although influential, the public goods model of alliances does not make accurate predictions about state size and defense spending. I find little evidence that larger states respond differently to changes in allied capability. Few alliances provide public goods in a way that leads to the collective action problem envisaged by Olson and Zeckhauser.

References

- Bolt, J, R Inklaar, H de Jong and JL van Zanden. 2018. “Maddison Project Database, Version 2018.” *Rebasing Maddison: new income comparisons and the shape of long-run economic development* .
- Brambor, Thomas, William Roberts Clark and Matt Golder. 2006. “Understanding interaction models: Improving empirical analyses.” *Political Analysis* 14(1):63–82.
- Carpenter, Bob, Andrew Gelman, Matt Hoffman, Daniel Lee, Ben Goodrich, Michael Betancourt, Michael A Brubaker, Jiqiang Guo, Peter Li and Allen Riddell. 2016. “Stan: A probabilistic programming language.” *Journal of Statistical Software* 20:1–37.
- DiGiuseppe, Matthew and Paul Poast. 2016. “Arms versus Democratic Allies.” *British Journal of Political Science* pp. 1–23.
- Gates, William R and Katsuaki L Terasawa. 1992. “Commitment, Threat Perceptions, and Expenditures in a Defense Alliance.” *International Studies Quarterly* 36(1):101–118.
- George, Justin and Todd Sandler. 2017. “Demand for military spending in NATO, 1968–2015: A spatial panel approach.” *European Journal of Political Economy* pp. 1–15.
- Gibler, Douglas M, Steven V Miller and Erin K Little. 2016. “An Analysis of the Militarized Interstate Dispute (MID) Dataset, 1816–2001.” *International Studies Quarterly* 60(4):719–730.
- Granger, Clive WJ and Paul Newbold. 1974. “Spurious Regressions in Econometrics.” *Journal of Econometrics* 2(2):111–120.

- Hainmueller, Jens, Jonathan Mummolo and Yiqing Xu. 2019. "How Much Should We Trust Estimates from Multiplicative Interaction Models?: Simple Tools to Improve Empirical Practice." *Political Analysis* .
- Kronmal, Richard A. 1993. "Spurious Correlation and the Fallacy of the Ratio Standard Revisited." *Journal of the Royal Statistical Society. Series A (Statistics in Society)* 156(3):379–92.
- Lanoszka, Alexander. 2015. "Do Allies Really Free Ride?" *Survival* 57(3):133–152.
- Leeds, Brett Ashley and Burcu Savun. 2007. "Terminating Alliances: Why Do States Abrogate Agreements?" *The Journal of Politics* 69(4):1118–1132.
- Leeds, Brett, Jeffrey Ritter, Sara Mitchell and Andrew Long. 2002. "Alliance Treaty Obligations and Provisions, 1815-1944." *International Interactions* 28(3):237–260.
- Morrow, James D. 1991. "Alliances and Asymmetry: An Alternative to the Capability Aggregation Model of Alliances." *American Journal of Political Science* 35(4):904–933.
- Olson, Mancur and Richard Zeckhauser. 1966. "An Economic Theory of Alliances." *The Review of Economics and Statistics* 48(3):266–279.
- Palmer, Glenn. 1990. "Alliance Politics and Issue Areas: Determinants of Defense Spending." *American Journal of Political Science* 34(1):190–211.
- Plümper, Thomas and Eric Neumayer. 2015. "Free-riding in alliances: Testing an old theory with a new method." *Conflict Management and Peace Science* 32(3):247–268.
- Rainey, Carlisle and Daniel K. Baissa. 2018. "When BLUE Is Not Best: Non-Normal Errors and the Linear Model." *Political Science Research & Methods* .
- Reiter, Dan, Allan C. Stam and Michael C. Horowitz. 2016. "A Revised Look at Interstate Wars, 1816–2007." *Journal of Conflict Resolution* 60(5):956–976.
- Sandler, Todd and John F Forbes. 1980. "Burden sharing, strategy, and the design of NATO." *Economic inquiry* 18(3):425–444.
- Sandler, Todd and Keith Hartley. 2001. "Economics of alliances: The lessons for collective action." *Journal of Economic Literature* 39(3):869–896.
- Sarkees, Meredith Reid and Frank Whelon Wayman. 2010. *Resort to war: a data guide to interstate, extra-state, intra-state, and non-state wars, 1816-2007*. Washington, DC: CQ Press.
- Singer, J David. 1988. "Reconstructing the correlates of war dataset on material capabilities of states, 1816–1985." *International Interactions* 14(2):115–132.